

Report on MLS Player Annual Earning Prediction Linear Regression Model

Introduction:

The high-salary phenomenon in the sports field attracts a lot of attention from the society. As a men's professional soccer league in the United States(US Soccer, 2021)MLS is a, and some soccer players have extremely high salaries.

By fitting a linear regression model on various statistics of MLS players and their salaries in 2018, I would like to answer my research question: What is the relationship between the annual earning of a MLS player and his performance in the year?

The importance of this research is that soccer players know their earnings relate to which statistics, so they can have a direction for improvement. Also, soccer clubs can better manage their budget to optimize their benefit, without choosing overvalued players.

Background/Literature:

Through a search on the UofT library, I found a paper about soccer clubs' choices in Spanish and English leagues from 1994 to 2004 that shows that win maximization better approximates the choices of soccer clubs than profit maximization(Garcia-del-Barrio & Szymanski, 2009). It shows the importance of my research question because we need to determine whether buying a player with a high salary really relates to victory. Also, a study analyzed the productivity and salary structure in MLS using the Gini coefficient and coefficient of variation. The researcher concluded that as salary increases, production declines(Coates et al., 2016). It verifies the inflation of salary among MLS players, which helps me anticipate the distribution of salary among MLS players: They will be highly skewed. It is important because to fit a linear regression model on earnings, I probably need to adjust my response variable.

Another study using positive assortative matching based on productivity showed that the high performance in MLS is related to wages(Scarfe et al., 2018). The study used data from previous years, and I expect I can reproduce and agree with the result by using a different approach. In addition, an analysis of the wage pattern in MLS is done using generalized least squares estimation. One of the most influential determinants they found is the number of games the player started with(Celik & Ince-Yenilmez, 2017). The approach they used is similar to mine, which means that I can include this variable as one of the variables in my analysis, although I have a different focus on predictors from them.

Data source:

I am using 2 datasets for data of their common MLS players' statistic in 2018. In my analysis, I will determine which predictors should be kept.

Main variables:

- **Total.Compensation:** Average annual guaranteed compensation of a MLS player.
- **GP:** Number of game a MLS player in.
- **GS:** Number of game a MLS player in when the game started.
- **G:** Number of goals scored by a MLS player.
- **A:** Number of on target scoring attempts of a MLS player.
- **SHTS:** Total scoring attempts of a MLS player.
- **SOG:** On target scoring attempts of a MLS player.
- **FS:** Number of fouls(unfair act) a MLS player suffered from.
- **OFF:** Number of times a MLS player is in an offside position, that is when he is nearer to his opponents's goal line than both the ball and the second last opponent(Rothschild, 2016).

The **Total.Compensation** data is scraped by Ben Jones:

<https://data.world/dataremixed/mls-player-salaries-2010-2018>

The original source of the data is: <https://mlsplayers.org/resources/salary-guide>

The **rest of main variables** are scraped by Joseph Mohr:

<https://www.kaggle.com/josephvm/major-league-soccer-dataset>

The original source of the data is: <https://www.mlssoccer.com/stats>

Sample Reliability, Accuracy and Sampling Bias

The original source of data is MLSPA and MLS. MLSPA is an organization that served as the collective bargaining representative for MLS players(MLSPA, 2021), while MLS is the official website of the league. Therefore, I anticipate the data is reliable and accurate. Data is directly scraped by from the source so the sampling bias that is most likely to occur is selection bias.

Generalization and Confounding Factors

After a natural join, I ended up with 690 different observations in my dataset. Generally, MLS has 776 players participating each year (Transfermarkt, 2021), so I expect that the result can be generalized to other MLS players. However, I failed to collect some confounding variables, and they may affect the generalizability of the result to a larger population. Possible confounding variables are the level of the soccer league, the budget of the team, popularity, age, and cultural differences.

Methods:

Salary was the average annual compensation of an MLS player, and it was the initial response in my model. *GP*, *GS*, *G*, *A*, *SHTS*, *SOG*, *FS*, *OFF* were different statistics of a player's performance.

Model Validation:

I randomly divided my dataset into a training set with 60%, and a validation dataset with 40% of the original dataset. Thus, I had enough data to generate a linear model, while the validation dataset contained similar data. I found my final model based on the training dataset and fitted it on the validation dataset. I tried to validate my model by checking if the model produced similar estimated coefficients, significant predictors, R^2 and adjusted R^2 . I also checked if new model violations and multicollinearity appeared. If the two models were similar, my model was validated.

Model Violations:

I first performed an exploratory data analysis(EDA) on the training set. Thus, potential issues such as assumption violations, multicollinearity, and influential points were anticipated. The initial model contained all possible predictors without transformations. I checked condition 1 by plotting the fitted value versus the response value. If points were randomly scattered around the identity function, I argued that condition 1 was satisfied. If no plot showed a non-linear relationship between predictors, I argued that condition 2 was satisfied. When conditions are satisfied, I plotted the residual plots for every variable and a QQ plot. I checked if the residuals are randomly scattered around 0 without a non-linear pattern, and if the QQ plot does not have a large deviation at the end. Combining all of my observations, I argued if assumptions were violated. To fix assumption violations while maintaining the interpretability, I used simple transformations such as logarithmic transformation referenced by box-cox transformation. I also used researches and context to argue the removal of problematic variables.

Variable Selection:

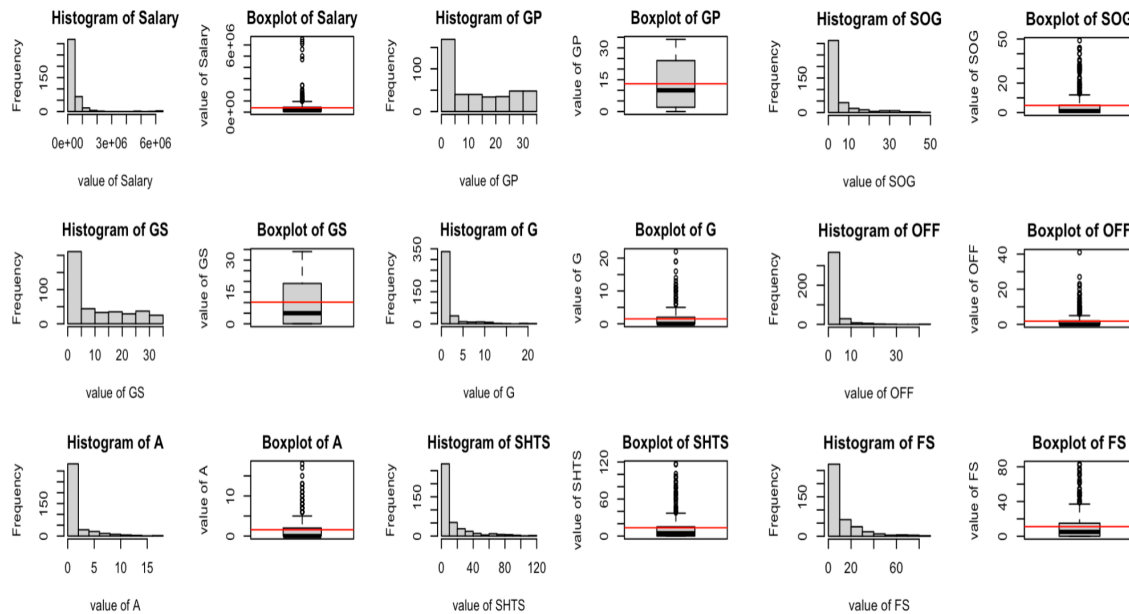
With a model that satisfied all assumptions, I checked the multicollinearity by calculating variance inflation factors(VIF) for each predictor. If the VIF exceeds 5, the predictor had severe multicollinearity. They could increase standard errors and make coefficients the wrong sign, so I removed some of them based on researches, context, and significance of variables. I made sure no new assumption violations appeared when problematic predictors were removed. I used stepwise selection based on Bayesian Information Criterion(BIC) because it accounted for the conditional nature of the model comparing to the forward and backward selection. I applied it for both sets of possible predictors with or without multicollinearity removed.

For the two resulting models, I checked assumptions and multicollinearity again to see if they were concerning. I chose a model with high R^2 , small BIC, and without assumption violation and multicollinearity as my final model. I used a partial F test to confirm that other predictors should not be in the model.

After that, I detected leverage points with leverage higher than the cut-off and outliers with standardized residuals higher than the cut-off. I also found influential points using the cook's distance, DEFFECTS, DFBETAs. They are problematic as they could influence the placement of the regression line. By EDA there was no measuring error, so no observation should be removed.

Results:

I found no missing information in my training dataset. Potential predictors all had many zero values, so I should be cautious when applying transformations. All variables were highly right-skewed, then the assumption of normality was probably violated, and transformations were needed. There were also many outliers for variables except for GP and GS, which means there might be problematic observations in the resulting model.



Distributions of Variables

For the initial model, it was enough to say conditions 1 and 2 were satisfied. The residual plots were all fine, but there was a downward trend in the residual versus fitted values plot, and the QQ plot went very upward at the end. Thus, I applied a log transformation on all variables. Since predictors had many zero values, I added 0.5 to the observations of all predictors. The new model satisfied conditions 1 and 2 well and showed a good QQ plot and residual plots. I concluded that this model satisfied the linear regression assumptions, which meant I could use the result from this model to make decisions.

I calculated VIF for all the predictors and found that $\log(GP)$, $\log(GS)$, $\log(SHTS)$, $\log(SOG)$, $\log(FS)$ had $VIF > 5$. From the correlation plot in the appendix, I knew that $\log(GP)$ and $\log(GS)$

were positively linearly related. Since GS, the number of games the player started, was previously shown to be a significant factor for wage (Celik & Ince-Yenilmez, 2017), I did not remove it. Instead, I removed GP, the number of games the player played, which had a similar definition as GS. Other predictors were not significant in the transformed model, so it was reasonable for me to remove them to resolve multicollinearity.

I applied stepwise selection methods based on BIC for both sets of transformed predictors with or without multicollinearity removed, and ended up with two potential models.

| | Model 1 | Model 2 |
|---------------------------|--|---|
| Initial Multicollinearity | Not Removed | Removed |
| Model | $\log(\text{Salary}) \sim \log(G) + \log(GS) + \log(GP)$ | $\log(\text{Salary}) \sim \log(G) + \log(GS)$ |
| Final Multicollinearity | $\log(GS), \log(GP)$ | None |
| BIC | -99.09 | -96.79 |
| R^2 | 0.318 | 0.3041 |
| Adjusted R^2 | 0.313 | 0.301 |
| Multicollinearity | None | None |

Comparing Models Produced by Automated Selection

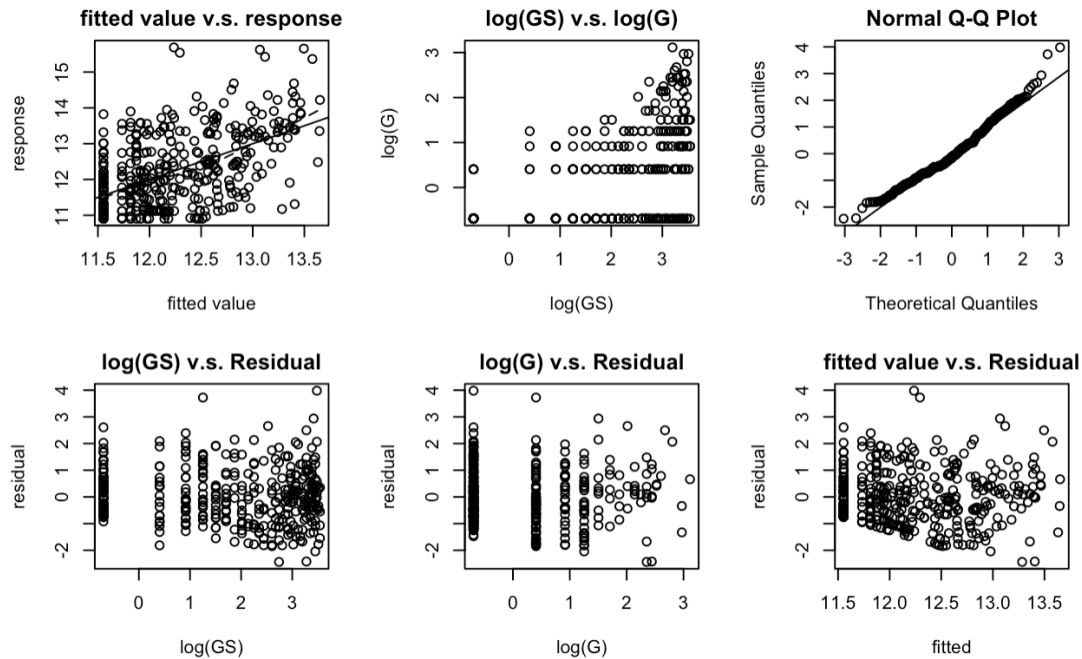
I performed partial F tests on transformed model versus Model 1, reduced model versus Model 2, and Model 1 versus Model 2. Only the last one showed a significant p-value, meaning that only $\log(G)$, $\log(GS)$, $\log(GP)$ should remain in the model.

Since Model 1 was Model 2 with an extra predictor, I compared their adjusted R^2 . Model 2's adjusted R^2 and R^2 were only 0.01 unit lower than Model 1.

Although Model 1 had a higher BIC, it had severe multicollinearity on $\log(GP)$ and $\log(GS)$. Overall, I argued that Model 2 was more preferable and was my final model.

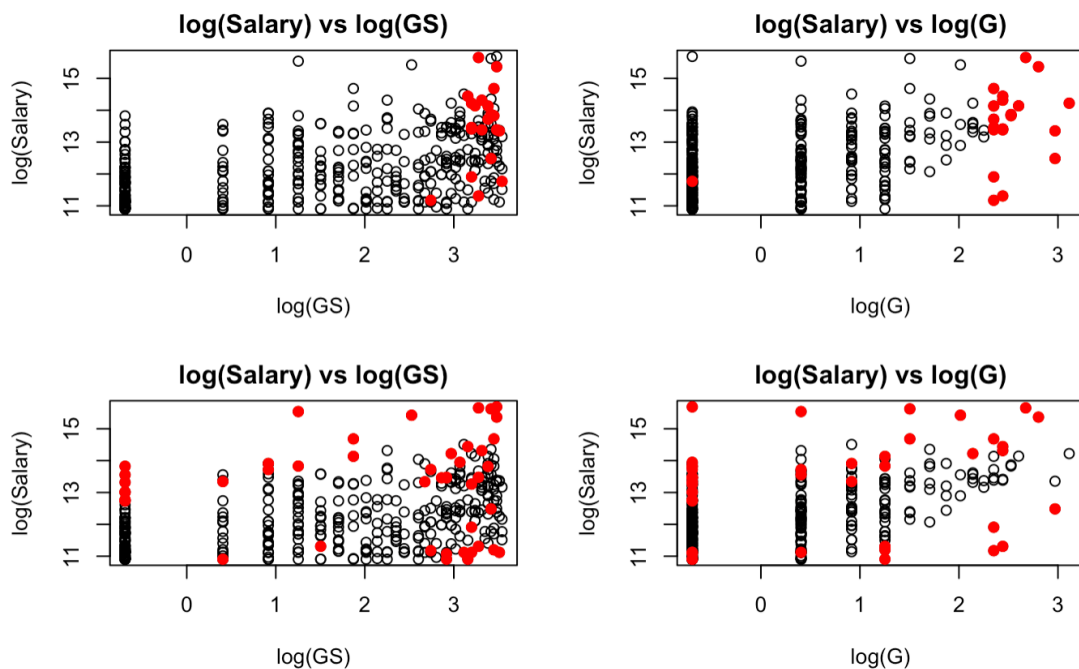
Goodness of Final Model:

Based on the plot of fitted value versus response, Model 2 satisfied condition 1. Although the graph of $\log(G)$ versus $\log(GS)$ had some fanning pattern, it might be due to the zeros in both variables. Considering the points were also scattered, it was enough to say condition 2 was satisfied. The QQ plot and residual plots looked fine, so I argued that the final model satisfies the assumptions of linear regression models.



Condition 1 & 2, Residual Plots and QQ Plot of Final Model

The model had no outliers, but some leverage points are shown in the first row of the following plots. There were also a few influential points, shown in the second row.



Leverage Points(1st Row), Influential Points(2nd Row)

| Final Model: $\log(\text{Salary}) \sim \log(\text{GS}) + \log(\text{G})$ | | | | | |
|--|-------------------------|------------|--------------|---------------------------|--------------|
| | Training Dataset | | | Validation Dataset | |
| | Estimate | Std. Error | Significance | Estimate | Significance |
| Intercept | 11.9 | 0.07 | <0.001 | 12 | <0.001 |
| $\log(\text{GS})$ | 0.16 | 0.03 | <0.001 | 0.12 | <0.05 |
| $\log(\text{G})$ | 0.38 | 0.05 | <0.001 | 0.33 | <0.001 |
| Residual standard error | 0.8736 | | | 0.9301 | |
| R^2 | 0.3041 | | | 0.219 | |
| Adjusted R^2 | 0.30 | | | 0.2133 | |
| Multicollinearity | None | | | None | |
| Leverage Points | 21 | | | 16 | |
| Outliers | None | | | None | |
| Influential Points | 46 | | | 26 | |

Comparing Training Model and Validation Model

Validation of Final Model:

For the model fitted on the validation dataset, all estimated coefficients were still significant and within 2 standard errors from the training model coefficients. It also had a similar residual standard error and number of leverage points and outliers. As shown in the Appendix, no new assumption violations, multicollinearity appeared in the validation model. There were fewer influential points which were reasonable since the dataset was smaller. Although the adjusted R^2 decreased by 0.9, most values corresponded to the training model. Overall, there was enough evidence to validate my model.

Discussion:

The final model was: $\log(\text{Salary}) \sim \log(\text{GS}) + \log(\text{G})$, with the estimated coefficients presented in the table above. GS means number of games the player started and G means number of goals the player scored.

For example, it showed that when the logarithm of the number of goals the player scored increased by 1, the average logarithmic of the annual earning of the player increased by 0.38. The result from the model was valuable because it suggested players to practice scoring which had a significant positive impact on their salary. Soccer clubs knew that players in the starting lineup should have a higher salary.

Limitations:


Although the data was reliable since they are from official websites, sampling bias might occur if observations were lost during the scraping process.

Since the model only had two predictors, it might be under-fitting, and some predictors in the true relationship might not be included. They could be confounding variables that were hard to obtain, such as players' popularity, or predictors eliminated by researches and context at the beginning. Missing them affects the generalizability of the model to a larger population. There were many leverage and influential points. They had a disproportionate impact on the estimated regression line and made it less accurate. The automated selection did not know the impact of these observations, and did not check assumptions during the selection process. Thus, it might not choose the optimal solution. Also, it might be biased because we did not create a hypothesis but searched for the most significant results.

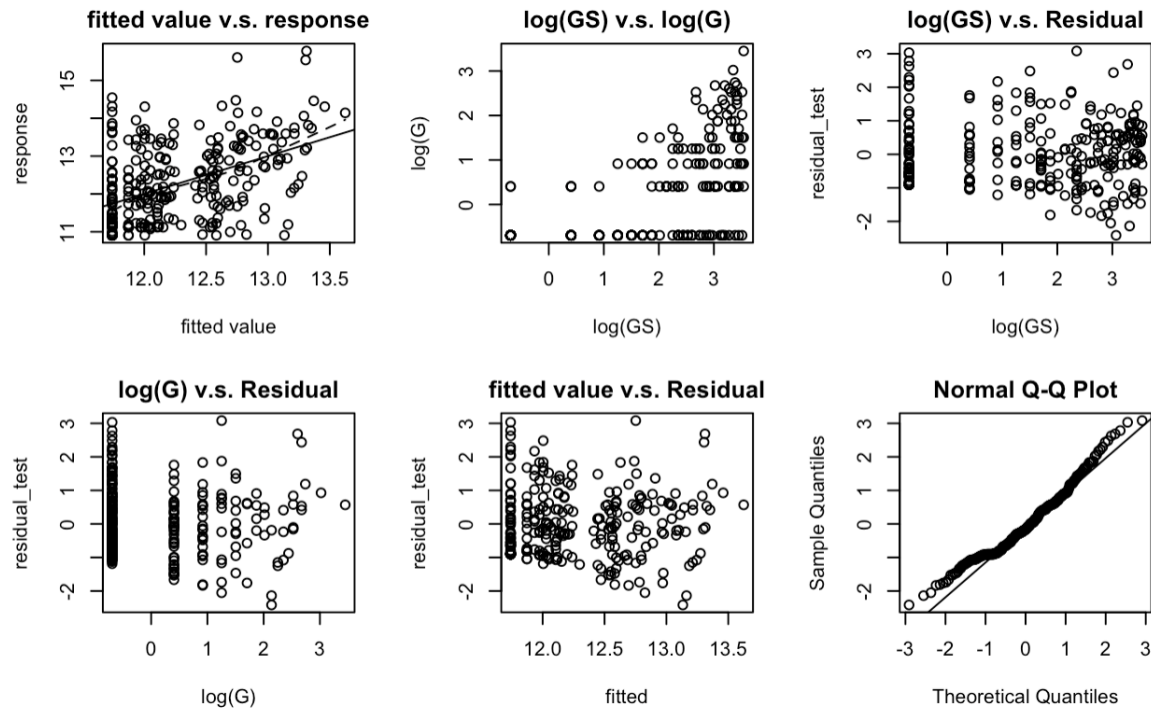
References

- Celik, O. B., & Ince-Yenilmez, M. (2017). Salary differences under the salary cap in Major League Soccer. *International Journal of Sports Science & Coaching*, 12(5), 623–634. <https://doi.org/10.1177/1747954117727809>
- Coates, D., Frick, B., & Jewell, T. (2016). Superstar salaries and soccer success. *Journal of Sports Economics*, 17(7), 716–735. <https://doi.org/10.1177/1527002514547297>
- Garcia-del-Barrio, P., & Szymanski, S. (2009). Goal! profit maximization versus win maximization in soccer. *Review of Industrial Organization*, 34(1), 45–68. <https://doi.org/10.1007/s11151-009-9203-6>
- Gastelum, A. (2021, September 21). *Ronaldo passes Messi as Forbes' top-paid soccer player*. Sports Illustrated. Retrieved October 20, 2021, from <https://www.si.com/soccer/2021/09/21/cristiano-ronaldo-highest-paid-soccer-player-forbes-messi-neymar>.
- Learn about the MLSPA*. MLS Players Association. (n.d.). Retrieved October 21, 2021, from <https://mlsplayers.org/about-us>.
- Major League Soccer 2021*. Transfermarkt. (n.d.). Retrieved October 20, 2021, from <https://www.transfermarkt.us/major-league-soccer/startseite/wettbewerb/MLS1>.
- Professional Leagues*. US Soccer. Sponsored by Volkswagen. (n.d.). Retrieved October 20, 2021, from <http://www.ussoccer.com/about/affiliates/professional-soccer>.
- Rothschild, T. (2016, January 21). *Soccer 101: Explaining the offside rule: Orlando City*. orlandocitysc. Retrieved October 20, 2021, from <https://www.orlandocitysc.com/news/soccer-101-explaining-offside-rule>.
- Scarfe, R., Singleton, C., & Telemo, P. (2018). Do high wage footballers play for high wage teams? *SSRN Electronic Journal*. <https://doi.org/10.2139/ssrn.3297644>

Appendix



Linearity Observed in Pairs of Predictors



Condition 1 & 2, Residual Plots, QQ plot for Final Model on Test Dataset